

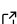
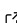
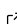
HiddenMarkovModels.jl: generic, fast and reliable latent variable modeling

Guillaume Dalle ¹

¹ EPFL (IdePHICS and INDY labs), Switzerland

DOI: [10.xxxxxx/draft](https://doi.org/10.xxxxxx/draft)

Software

- [Review](#) 
- [Repository](#) 
- [Archive](#) 

Editor: [Mehmet Hakan Satman](#) 

Submitted: 12 September 2023

Published: unpublished

License

Authors of papers retain copyright and release the work under a Creative Commons Attribution 4.0 International License ([CC BY 4.0](#)).

Summary

Hidden Markov Models (or HMMs) are a very popular statistical framework, with numerous applications ranging from speech recognition to bioinformatics. They model a sequence of observations Y_1, \dots, Y_T by assuming the existence of a hidden sequence of states X_1, \dots, X_T . The distribution of a state X_t can only depend on the previous state X_{t-1} , and the distribution of an observation Y_t can only depend on the current state X_t . This is a very versatile and practical set of assumptions: see Rabiner (1989) for an introduction and Cappé et al. (2005) for a book-length treatment.

Given a sequence of observations and a parametric family of HMMs \mathbb{P}_θ , there are several problems one can face. In generic graphical models, these problems are often intractable, but HMMs have a tree-like structure that yields exact solution procedures with polynomial complexity.

Problem	Algorithm
Observation sequence likelihood $\mathbb{P}_\theta(Y_{1:T})$	Forward
State marginals $\mathbb{P}_\theta(X_t Y_{1:T})$	Forward-backward
Best state sequence $\arg\max_{X_{1:T}} \mathbb{P}_\theta(X_{1:T} Y_{1:T})$	Viterbi
Maximum likelihood parameter $\arg\max_\theta \mathbb{P}_\theta(Y_{1:T})$	Baum-Welch

The package HiddenMarkovModels.jl leverages the Julia language (Bezanson et al., 2017) to implement those algorithms in a *generic, fast and reliable* way.

Statement of need

The initial motivation for HiddenMarkovModels.jl was an application of HMMs to reliability analysis for the French railway company SNCF (Dalle, 2022). In this industrial use case, the observations were marked temporal point processes (sequences of timed events with structured metadata) generated by condition monitoring systems.

Unfortunately, the major implementations of HMMs we surveyed (in Julia and Python) all expect the observations to be generated by a *predefined set of distributions*. In Julia, the reference package HMMBase.jl (Mouchet, 2023) requires compliance with the Distributions.jl (Besançon et al., 2021) interface, which precludes anything not scalar- or array-valued. In Python, hmmlearn (hmmlearn, 2023) only offers Gaussians, mixtures of Gaussians and a few discrete distributions (categorical, multinomial, Poisson). Meanwhile, pomegranate (Schreiber, 2018) has a wider set of available distributions, but it doesn't allow for easy extension by the user.

32 Focusing on Julia specifically, other downsides of HMMBase.jl include the lack of support for
33 multiple observation sequences, and the mandatory use of 64-bit floating point numbers. Two
34 other packages provide functionalities that HMMBase.jl lacks: HMMGradients.jl ([Antonello,
35 2021](#)) contains a differentiable loglikelihood function, while MarkovModels.jl ([Ondel et al.,
36 2021](#)) focuses on GPU acceleration. Unfortunately, all three have mutually incompatible APIs.

37 Package design

38 HiddenMarkovModels.jl was designed to overcome the limitations mentioned above, with the
39 following guiding principles in mind.

40 It is *generic*. Observations can be arbitrary objects, and the associated distributions only need
41 to implement two methods: a loglikelihood `logdensityof(dist, x)` and a sampler `rand(rng,
42 x)`. The extendable `AbstractHMM` interface allows incorporating features such as priors or
43 structured transition matrices. Number types are not restricted, and automatic differentiation
44 of the sequence loglikelihood ([Qin et al., 2000](#)) is supported both in forward and reverse mode.

45 It is *fast*. Julia's blend of multiple dispatch and just-in-time compilation delivers satisfactory
46 speed even when working with arbitrary observations. Inference routines rely on BLAS calls for
47 linear algebra, and exploit multithreading to process sequences in parallel.





48 It is *reliable*. The package is thoroughly tested and documented, with an extensive API
49 reference and accessible tutorials. Special care was given to code quality, type stability, and
50 compatibility checks with various downstream packages.

51 As a consequence, it is also *limited in scope*. It centers around CPU efficiency, and remains
52 untested on GPU. Its primary target is small- to medium-sized HMMs (a few tens of states),
53 mostly because memory requirements scale quadratically for the chosen storage mode. Finally,
54 it does not perform probability computations in the logarithmic domain, but instead uses the
55 scaling trick from Rabiner ([1989](#)) with a clever variation borrowed from HMMBase.jl. Thus,
56 its numerical stability might be lower than that of `hmmlearn` or `pomegranate` in challenging
57 instances. However, thanks to unrestricted number types, users are free to bring in third-party
58 packages like `LogarithmicNumbers.jl` ([Rowley, 2023](#)) for additional precision.

59 Benchmarks

60 We compare `HiddenMarkovModels.jl` (abbreviated to `HMMs.jl`), `HMMBase.jl`, `hmmlearn` and
61 `pomegranate` on a test case with multivariate Gaussian observations. The relevant parameters
62 are the number of states N , the sequence duration T , the observation dimension D , the number
63 of sequences K and the number of Baum-Welch iterations I . This benchmarking code is run
64 automatically by GitHub Actions upon each package release, but here we ran it on a local
65 Linux system with more cores. The numbers of Julia, OpenBLAS (for NumPy) and MKL
66 (for PyTorch) threads were all set to 6, although it is hard to compare all libraries fairly
67 when it comes to parallel computing. See the package documentation for more details on the
68 benchmarks.





69 In a low-dimensional scenario, `HiddenMarkovModels.jl` runs substantially faster than its
70 predecessor `HMMBase.jl`, even though their algorithms are mathematically identical. We note
71 that performance is less convincing for the Viterbi algorithm, in which linear operations are
72 replaced by max-plus operations (less easy to optimize). `pomegranate` is excluded because it
73 does not support multiple sequences.



74

75

In a high-dimensional scenario, HiddenMarkovModels.jl scales favorably compared to both Python alternatives, even though the latter make use of highly optimized C++ backends. This can partly be explained by the absence of logarithmic computations. HMMBase.jl is excluded because it does not support multiple sequences.



80

81

Conclusion

82

HiddenMarkovModels.jl fills a longstanding gap in the Julia package ecosystem, and might even prove interesting for Python users. Future research directions include the implementation

84

85 of Input-Output HMMs (Bengio & Frasconi, 1994) as well as other estimation methods
86 (gradient-based or spectral).

87 Acknowledgements

88 A special thank you goes to Maxime Mouchet and Jacob Schreiber, the developers of
89 HMMBase.jl and pomegranate respectively, for their help and advice. In particular, Maxime
90 agreed to designate HiddenMarkovModels.jl as the official successor to HMMBase.jl, which I
91 am very grateful for.

92 References

- 93 Antonello, N. (2021). *HMMGradients.jl: Enables computing the gradient of the parameters of*
94 *Hidden Markov Models (HMMs)*. Zenodo. <https://doi.org/10.5281/zenodo.4454565>
- 95 Bengio, Y., & Frasconi, P. (1994). An Input Output HMM Architecture. *Advances in Neural*
96 *Information Processing Systems*, 7. [https://proceedings.neurips.cc/paper/1994/hash/](https://proceedings.neurips.cc/paper/1994/hash/8065d07da4a77621450aa84fee5656d9-Abstract.html)
97 [8065d07da4a77621450aa84fee5656d9-Abstract.html](https://proceedings.neurips.cc/paper/1994/hash/8065d07da4a77621450aa84fee5656d9-Abstract.html)
- 98 Besançon, M., Papamarkou, T., Anthoff, D., Arslan, A., Byrne, S., Lin, D., & Pearson, J. (2021).
99 Distributions.jl: Definition and Modeling of Probability Distributions in the JuliaStats
100 Ecosystem. *Journal of Statistical Software*, 98, 1–30. <https://doi.org/10.18637/jss.v098.i16>
- 101 Bezanson, J., Edelman, A., Karpinski, S., & Shah, V. B. (2017). Julia: A Fresh Approach to
102 Numerical Computing. *SIAM Review*, 59(1), 65–98. <https://doi.org/10.1137/141000671>
- 103 Cappé, O., Moulines, E., & Rydén, T. (2005). *Inference in Hidden Markov Models*. Springer
104 New York. <https://doi.org/10.1007/0-387-28982-8>
- 105 Dalle, G. (2022). *Machine learning and combinatorial optimization algorithms, with applications*
106 *to railway planning* [PhD thesis, École des Ponts ParisTech]. [https://www.theses.fr/](https://www.theses.fr/2022ENPC0047)
107 [2022ENPC0047](https://www.theses.fr/2022ENPC0047)
- 108 hmmlearn. (2023). *Hmmlearn: Hidden Markov Models in Python, with scikit-learn like API*.
109 hmmlearn. <https://github.com/hmmlearn/hmmlearn>
- 110 Mouchet, M. (2023). *HMMBase.jl: Hidden Markov Models for Julia*. [https://github.com/](https://github.com/maxmouchet/HMMBase.jl)
111 [maxmouchet/HMMBase.jl](https://github.com/maxmouchet/HMMBase.jl)
- 112 Ondel, L., Lam-Yee-Mui, L.-M., Kocour, M., Filippo, C., & Lukás Burget, C. (2021). *GPU-*
113 *Accelerated Forward-Backward Algorithm with Application to Lattice-Free MMI*. [https://](https://hal.science/hal-03434552)
114 hal.science/hal-03434552
- 115 Qin, F., Auerbach, A., & Sachs, F. (2000). A Direct Optimization Approach to Hidden
116 Markov Modeling for Single Channel Kinetics. *Biophysical Journal*, 79(4), 1915–1927.
117 [https://doi.org/10.1016/S0006-3495\(00\)76441-1](https://doi.org/10.1016/S0006-3495(00)76441-1)
- 118 Rabiner, L. R. (1989). A tutorial on hidden Markov models and selected applications in speech
119 recognition. *Proceedings of the IEEE*, 77(2), 257–286. <https://doi.org/cswph2>
- 120 Rowley, C. (2023). *LogarithmicNumbers.jl: A logarithmic number system for Julia*. [https://](https://github.com/cjdoris/LogarithmicNumbers.jl)
121 github.com/cjdoris/LogarithmicNumbers.jl
- 122 Schreiber, J. (2018). Pomegranate: Fast and Flexible Probabilistic Modeling in Python. *Journal*
123 *of Machine Learning Research*, 18(164), 1–6. <http://jmlr.org/papers/v18/17-636.html>