

# Imbalance: A comprehensive multi-interface Julia toolbox to address class imbalance

Essam Wisam <sup>1\*</sup> and Anthony Blaom <sup>2\*</sup>

<sup>1</sup> Cairo University, Egypt <sup>2</sup> University of Auckland, New Zealand \* These authors contributed equally.

DOI: [10.xxxxxx/draft](https://doi.org/10.xxxxxx/draft)

## Software

- [Review](#) 
- [Repository](#) 
- [Archive](#) 

Editor: [Mehmet Hakan Satman](#) 

Submitted: 17 October 2023

Published: unpublished

## License

Authors of papers retain copyright and release the work under a Creative Commons Attribution 4.0 International License ([CC BY 4.0](#)).

## Summary

Given a set of observations that each belong to a certain class, supervised classification aims to learn a classification model that can predict the class of a new, unlabeled observation (Cunningham et al., 2008). This modeling process finds extensive application in real-life scenarios, including but not limited to medical diagnostics, recommendation systems, credit scoring, and sentiment analysis.

In various real-world scenarios where supervised classification is employed, such as those pertaining to the detection of particular conditions like fraud, faults, pollution, or rare diseases, a severe discrepancy between the number of observations in each class can occur. This is known as class imbalance. This poses a problem if assumptions inherent in the classification model imply hindered performance when the model is trained on imbalanced data as is commonly the case (Ali et al., 2015). Two prevalent strategies for mitigating class imbalance, when it poses a problem to the classification model, involve either increasing the representation of less frequently occurring classes through oversampling or reducing instances of more frequently occurring classes through undersampling. It may be also possible to achieve even greater performance by combining both approaches (Zeng et al., 2016) or by resampling the data multiple times and training the classification model on each resampled dataset to form an ensemble model that aggregates results from different model instances (Liu et al., 2009).

## Imbalance.jl

In this work, we present, `Imbalance.jl`, a software toolbox implemented in the Julia programming language that offers over 10 well-established techniques that help address the class imbalance issue. Additionally, we present a companion package, `MLJBalancing.jl`, which: (i) facilitates the integration of resampling methods with classification models via the `BalancedModel` construct, to create a seamless machine learning pipeline that behaves like a single unified model; and (ii) implements a general version of the `EasyEnsemble` algorithm presented in (Liu et al., 2009). The set of resampling techniques implemented in `Imbalance.jl` and `MLJBalancing.jl` are shown in ???. Although no combination resampling techniques are explicitly presented, they are easy to form using the `BalancedModel` wrapper found in `MLJBalancing.jl`.

The toolbox offers a pure functional interface for each method implemented. For example, `SMOTE` can be used in the following fashion:

```
Xover, yover = smote(X, y)
```

Here `Xover`, `yover` are `X`, `y` after oversampling.

A `ratios` hyperparameter or similar is always present to control the degree of oversampling or undersampling to be done for each class. All hyperparameters for a resampling method have default values that can be overridden.

**Table 1:** Resampling techniques implemented in Imbalance.jl and MLJBalancing.jl.

Technique	Type	Supported Data Types
BalancedBaggingClassifier	Ensemble	Continuous and/or nominal
Borderline SMOTE1	Oversampling	Continuous
Cluster Undersampler	Undersampling	Continuous
Edited Nearest Neighbors Undersampler	Undersampling	Continuous
Random Oversampler	Oversampling	Continuous and/or nominal
Random Undersampler	Undersampling	Continuous and/or nominal
Random Walk Oversampler	Oversampling	Continuous and/or nominal
ROSE	Oversampling	Continuous
SMOTE	Oversampling	Continuous
SMOTE-N	Oversampling	Nominal
SMOTE-NC	Oversampling	Continuous and nominal
Tomek Links Undersampler	Undersampling	Continuous

## Statement of Need

A substantial body of literature in the field of machine learning and statistics is devoted to addressing the class imbalance issue. This predicament has often been aptly labeled the “curse of class imbalance,” as noted in (Picek et al., 2018) and (Kubát & Matwin, 1997) which follows from the pervasive nature of the issue across diverse real-world applications and its pronounced severity; a classifier may incur an extraordinarily large performance penalty in response to training on imbalanced data.

The literature encompasses a myriad of oversampling and undersampling techniques to approach the class imbalance issue. These include SMOTE (Chawla et al., 2002) which operates by generating synthetic examples along the lines joining existing points, SMOTE-N and SMOTE-NC (Chawla et al., 2002) which are variants of SMOTE that can deal with categorical data. The sheer number of SMOTE variants makes them a body of literature on their own. Notably, the most widely cited variant of SMOTE is BorderlineSMOTE (Han et al., 2005). Other well-established oversampling techniques include RWO (Zhang & Li, 2014) and ROSE (Menardi & Torelli, 2012). On the other hand, the literature also encompasses many undersampling techniques such as cluster undersampling (Lin et al., 2016) and condensed nearest neighbors (Hart, 1968). Furthermore, methods that combine oversampling and undersampling (Zeng et al., 2016) or resampling with ensemble learning (Liu et al., 2009) are also present.

The existence of a toolbox with techniques that harness this wealth of research is necessary for the development of novel approaches to the class imbalance problem and for machine learning research in general. Aside from addressing class imbalance in a general machine learning research setting, the toolbox can help in class imbalance research settings by making it possible to juxtapose different methods, compose them together, or form variants of them without having to reimplement them from scratch. In popular programming languages, such as Python, a variety of such toolboxes already exist, such as imbalanced-learn (Lemaître et al., 2016) and SMOTE-variants (Kovács, 2019). Meanwhile, Julia, a well known programming language with over 40M downloads (Tuychiev, 2023), has been lacking a similar toolbox to address the class imbalance issue in general multi-class, heterogeneous data settings. This has served as the primary motivation for the creation of the Imbalance.jl toolbox.

## Imbalance.jl Design Principles

The toolbox implementation follows a specific set of design principles in terms of the implemented techniques, interface support, developer experience and testing, and user experience.

## 72 Implemented Techniques

- 73     ▪ Should support all four major types of resampling approaches
- 74     ▪ Should be generally compatible with multi-class settings
- 75     ▪ Should offer solutions to heterogeneous data settings (continuous and nominal data)
- 76     ▪ When possible, preference should be given to techniques that are more common in the
- 77         literature or industry

78 Methods implemented in the `Imbalance.jl` toolbox indeed meet all aforementioned design  
79 principles for the implemented techniques. The one-vs-rest scheme as proposed in (Fernández  
80 et al., 2013) was used to generalize the technique to multi-class when needed.

## 81 Interface Support

- 82     ▪ Should support both matrix and table inputs
- 83     ▪ Target variable may or may not be given as a separate column
- 84     ▪ Should expose a pure functional implementation, but also support popular Julia machine
- 85         learning interfaces
- 86     ▪ Should be possible to wrap an arbitrary number of resampler models with an MLJ model
- 87         to behave as a unified model using `MLJBalancing`

88 Methods implemented in the `Imbalance.jl` toolbox meet all the interface design principles  
89 above. It particularly implements the `MLJ` and `TableTransforms` interface for each method.  
90 `BalancedModel` from `MLJBalancing.jl` also allows fusing an arbitrary number of resampling  
91 models and a classifier together to behave as one unified model.

## 92 Developer Experience and Testing

- 93     ▪ Should document all functions, including internal ones
- 94     ▪ Comments should be included to justify or simplify written implementations when needed
- 95     ▪ Features commonly used by multiple resampling techniques should be implemented in a
- 96         single function and reused
- 97     ▪ Functions should be implemented in smaller units to aid in testing
- 98     ▪ Testing coverage should be maximized; even the most basic functions should be tested
- 99     ▪ There should exist a developer guide to encourage and guide contribution

100 This set of design principles is also satisfied by `Imbalance.jl`. Implemented techniques are  
101 tested by testing smaller units that form the technique. End-to-end tests are performed for  
102 each technique by testing properties and characteristics of the technique or by using the  
103 `imbalanced-learn` toolbox from Python and comparing outputs.

## 104 User Experience

- 105     ▪ Functional documentation should be comprehensive and clear
- 106     ▪ Examples (with shown output) that work after copy-pasting should accompany each
- 107         method
- 108     ▪ An illustrative visual example that presents a plot or animation should preferably accom-
- 109         pany each method
- 110     ▪ A practical example that uses the method with real data should preferably accompany
- 111         each method
- 112     ▪ Users should preferably be able to easily run the illustrative or practical examples (e.g.,
- 113         via Google Colab)
- 114     ▪ If an implemented method lacks an online explanation, an article that explains the
- 115         method after it is implemented should be preferably written

116 The `Imbalance.jl` documentation indeed satisfies this set of design principles. Methods are  
117 associated with examples that can be copy-pasted, examples that demonstrate the operation of

the technique visually, and possibly, examples that use it with a real-world dataset to improve the performance of a classification model.

## Author Contributions

Design: E. Wisam, A. Blaom. Implementation, tests and documentation: E. Wisam. Code and documentation review: A. Blaom.

## References

- Ali, A., Shamsuddin, S. M. Hj., & Ralescu, A. L. (2015). Classification with class imbalance problem: A review. *Soft Computing Models in Industrial and Environmental Applications*. <https://api.semanticscholar.org/CorpusID:26644563>
- Chawla, N., Bowyer, K., Hall, L. O., & Kegelmeyer, W. P. (2002). SMOTE: Synthetic minority over-sampling technique. *ArXiv, abs/1106.1813*. <https://api.semanticscholar.org/CorpusID:1554582>
- Cunningham, P., Cord, M., & Delany, S. J. (2008). Supervised learning. In M. Cord & P. Cunningham (Eds.), *Machine learning techniques for multimedia: Case studies on organization and retrieval* (pp. 21–49). Springer Berlin Heidelberg. [https://doi.org/10.1007/978-3-540-75171-7\\_2](https://doi.org/10.1007/978-3-540-75171-7_2)
- Fernández, A., López, V., Galar, M., Jesús, M. J. del, & Herrera, F. (2013). Analysing the classification of imbalanced data-sets with multiple classes: Binarization techniques and ad-hoc approaches. *Knowl. Based Syst.*, 42, 97–110. <https://api.semanticscholar.org/CorpusID:131286>
- Han, H., Wang, W., & Mao, B. (2005). Borderline-SMOTE: A new over-sampling method in imbalanced data sets learning. *International Conference on Intelligent Computing*. <https://api.semanticscholar.org/CorpusID:12126950>
- Hart, P. E. (1968). The condensed nearest neighbor rule (corresp.). *IEEE Trans. Inf. Theory*, 14, 515–516. <https://api.semanticscholar.org/CorpusID:206729609>
- Kovács, G. (2019). Smote-variants: A python implementation of 85 minority oversampling techniques. *Neurocomputing*, 366, 352–354. <https://doi.org/https://doi.org/10.1016/j.neucom.2019.06.100>
- Kubát, M., & Matwin, S. (1997). Addressing the curse of imbalanced training sets: One-sided selection. *International Conference on Machine Learning*. <https://api.semanticscholar.org/CorpusID:18370956>
- Lemaître, G., Nogueira, F., & Aridas, C. K. (2016). Imbalanced-learn: A python toolbox to tackle the curse of imbalanced datasets in machine learning. *ArXiv, abs/1609.06570*. <https://api.semanticscholar.org/CorpusID:1426815>
- Lin, W.-C., Tsai, C.-F., Hu, Y.-H., & Jhang, J.-S. (2016). Clustering-based undersampling in class-imbalanced data. *Inf. Sci.*, 409, 17–26. <https://api.semanticscholar.org/CorpusID:424467>
- Liu, X.-Y., Wu, J., & Zhou, Z.-H. (2009). Exploratory undersampling for class-imbalance learning. *IEEE Transactions on Systems, Man, and Cybernetics, Part B (Cybernetics)*, 39, 539–550. <https://api.semanticscholar.org/CorpusID:62808464>
- Menardi, G., & Torelli, N. (2012). Training and assessing classification rules with imbalanced data. *Data Mining and Knowledge Discovery*, 28, 92–122. <https://api.semanticscholar.org/CorpusID:18164904>
- Picek, S., Heuser, A., Jović, A., Bhasin, S., & Regazzoni, F. (2018). The curse of class imbalance and conflicting metrics with machine learning for side-channel evaluations. *IACR*

- 163 *Trans. Cryptogr. Hardw. Embed. Syst.*, 2019, 209–237. [https://api.semanticscholar.org/](https://api.semanticscholar.org/CorpusID:44136202)  
164 [CorpusID:44136202](https://api.semanticscholar.org/CorpusID:44136202)
- 165 Tuychiev, B. (2023). *The rise of julia*. <https://www.datacamp.com/blog/the-rise-of-julia-is-it-worth-learning>
- 166 Zeng, M., Zou, B., Wei, F., Liu, X., & Wang, L. (2016). Effective prediction of three common  
167 diseases by combining SMOTE with totem links technique for imbalanced medical data.  
168 *2016 IEEE International Conference of Online Analysis and Computing Science (ICOACS)*,  
169 225–228. <https://api.semanticscholar.org/CorpusID:25184489>
- 170 Zhang, H., & Li, M. (2014). RWO-sampling: A random walk over-sampling approach to  
171 imbalanced data classification. *Inf. Fusion*, 20, 99–116. [https://api.semanticscholar.org/](https://api.semanticscholar.org/CorpusID:205432428)  
172 [CorpusID:205432428](https://api.semanticscholar.org/CorpusID:205432428)

DRAFT