

# Taweret: a Python package for Bayesian model mixing

K. Ingles<sup>1\*</sup>, D. Liyanage<sup>2\*</sup>, A. C. Sempowski<sup>3\*</sup>, and J. C. Yannotty<sup>4\*</sup>

<sup>1</sup> Illinois Center for Advanced Study of the Universe & Department of Physics, University of Illinois Urbana-Champaign, USA <sup>2</sup> Department of Physics, The Ohio State University, USA <sup>3</sup> Department of Physics and Astronomy & Institute of Nuclear and Particle Physics, Ohio University, USA <sup>4</sup> Department of Statistics, The Ohio State University, USA ¶ Corresponding author \* These authors contributed equally.

DOI: [10.xxxxxx/draft](https://doi.org/10.xxxxxx/draft)

## Software

- Review
- Repository
- Archive

Editor: Rachel Kurchin

## Reviewers:

- @julienmalard
- @gchure

Submitted: 31 October 2023

Published: unpublished

## License

Authors of papers retain copyright and release the work under a Creative Commons Attribution 4.0 International License ([CC BY 4.0](https://creativecommons.org/licenses/by/4.0/)).

## Summary

Uncertainty quantification using Bayesian methods is a growing area of research. Bayesian model mixing (BMM) is a recent development which combines the predictions from multiple models such that each model's best qualities are preserved in the final result. Practical tools and analysis suites that facilitate such methods are therefore needed. Taweret introduces BMM to existing Bayesian uncertainty quantification efforts. Currently Taweret contains three individual Bayesian model mixing techniques, each pertaining to a different type of problem structure; we encourage the future inclusion of user-developed mixing methods. Taweret's first use case is in nuclear physics, but the package has been structured such that it should be adaptable to any research engaged in model comparison or model mixing.

## Statement of need

In physics applications, multiple models with different physics assumptions can be used to describe an underlying system of interest. It is usually the case that each model has varying fidelity to the observed process across the input domain. Though each model may have similar predictive accuracy on average, the fidelity of the approximation across a subset of the domain may differ drastically for each of the models under consideration. In such cases, inference and prediction based on a single model may be unreliable. One strategy for improving accuracy is to combine, or “mix”, the predictions from each model using a linear combination or weighted average with input-dependent weights. This approach is intended to improve reliability of inference and prediction and properly quantify model uncertainties. When operating under a Bayesian framework, this technique is referred to as Bayesian model mixing (BMM). In general, model mixing techniques are designed to combine the individual mean predictions or density estimates from the  $K$  models under consideration. For example, *mean-mixing* techniques predict the underlying system by

$$E[Y | x] = \sum_{k=1}^K w_k(x) f_k(x),$$

where  $E[Y | x]$  denotes the mean of  $Y$  given the vector of input parameters  $x$ ,  $f_k(x)$  is the mean prediction under the  $k^{\text{th}}$  model  $\mathcal{M}_k$ , and  $w_k(x)$  is the corresponding weight function. The *density-mixing* approach estimates the underlying predictive density by

$$p(Y_0 | x_0, Y) = \sum_{k=1}^K w_k(x_0) p(Y_0 | x_0, Y, \mathcal{M}_k),$$

where  $p(Y_0 | x_0, Y, \mathcal{M}_k)$  represents the predictive density of a future observation  $Y_0$  with respect to the  $k^{\text{th}}$  model  $\mathcal{M}_k$  at a new input  $x_0$ . In either BMM setup, a key challenge is

37 defining  $w_k(x)$ —the functional relationship between the inputs and the weights.

38 This work introduces Taweret, a Python package for Bayesian model mixing that includes  
 39 three novel approaches for combining models, each of which defines the weight function in a  
 40 unique way (see Table 1 for a comparison of each method). This package has been developed  
 41 as an integral piece of the Bayesian Analysis of Nuclear Dynamics (BAND) collaboration's  
 42 software. BAND is a multi-institutional effort to build a cyber-infrastructure framework for  
 43 use in the nuclear physics community (Beyer et al., 2023; Phillips et al., 2021). The software  
 44 is designed to lower the barrier for researchers to employ uncertainty quantification in their  
 45 experiments, and to integrate, as best as possible, with the community's current standards  
 46 concerning coding style (pep8). Bayesian model mixing is one of BAND's four central pillars  
 47 in this framework (the others being emulation, calibration, and experimental design).

48 In addition to this need, we are aware of several other fields outside of physics that use  
 49 techniques such as model stacking and Bayesian model averaging (BMA) (Fragoso et al.,  
 50 2018), e.g., statistics (Yao et al., 2018, 2022), meteorology (Slougher et al., 2007), and  
 51 neuroscience (FitzGerald et al., 2014). It is expected that the Bayesian model mixing methods  
 52 presented in Taweret can also be applied to use cases within these fields. Statisticians have  
 53 developed several versatile BMA/stacking packages, e.g. (Raftery et al., 2022; Vehtari et al.,  
 54 2017). However, the only BMM-based package available is SAMBA—a BAND collaboration  
 55 effort that was developed for testing BMM methods on a toy model (Semposki et al., 2022a).  
 56 Taweret's increased functionality, user-friendly structure, and diverse selection of mixing  
 57 methods make it a marked improvement over SAMBA.

## 58 Structure

### 59 Overview of methods

**Table 1:** A summary of the three BMM approaches currently implemented in Taweret. Note that  $K \geq 2$ . Following the method name and the type of mixing model, the *Number of inputs* column details the dimensions of the parameter which the mixing weights depend on (e.g., in heavy-ion collisions this is the centrality bin); the *Number of outputs* details how many observables the models themselves can have to compute the model likelihood (e.g., in heavy-ion collisions this can include charge multiplicities, transverse momentum distributions, transverse momentum fluctuations, etc.); the *Number of models* column details how many models the mixing method can combine, and the *Weight functions* column describes the available parameterization of how the mixing weights depend on the input parameter.

Method	Type	Number of inputs	Number of outputs	Number of models	Weight functions
Bivariate linear mixing	Mean & Density	1	$\geq 1$	2	<ul style="list-style-type: none"> <li>Step,</li> <li>Sigmoid,</li> <li>Asymmetric 2-step Precision weighting</li> </ul>
Multivariate mixing	Mean	1	1	$K$	
BART mixing	Mean	$\geq 1$	1	$K$	Regression trees

### 60 Bivariate linear mixing

61 The full description of this mixing method and several of its applications in relativistic heavy-ion  
 62 collision physics can be found in the Ph.D. thesis of D. Liyanage (Liyanage, 2023). The bivariate  
 63 linear mixing method can mix two models either using a density-mixing or a mean-mixing  
 64 strategy. Currently, this is the only mixing method in Taweret that can also calibrate the  
 65 models while mixing. It allows the user to choose among the following mixing functions:

- step:  $\Theta(\beta_0 - x)$
- sigmoid:  $\exp[(x - \beta_0)/\beta_1]$
- asymmetric 2-step:  $\alpha\Theta(\beta_0 - x) + (1 - \alpha)\Theta(\beta_1 - x)$ .

Here  $\Theta$  denotes the Heaviside step function,  $\beta_0$  and  $\beta_1$  determine the shape of the weight function and are inferred from the experimental data, and  $x$  is the model input parameter (which is expected to be 1-dimensional for this mixing method).

## Multivariate model mixing

Another Bayesian model mixing method incorporated into Taweret was originally published in (Sempowski et al., 2022a), and was the focus of the BMM Python package SAMBA (Sempowski et al., 2022b). It can be described as combining models by weighting each of them by their precision, defined as the inverse of their respective variances. The posterior predictive distribution (PPD) of the mixed model is a Gaussian and can be expressed as

$$\mathcal{M}_{\dagger} \sim \mathcal{N}(f_{\dagger}, Z_P^{-1}) : \quad f_{\dagger} = \frac{1}{Z_P} \sum_{k=1}^K \frac{1}{\sigma_k^2} f_k, \quad Z_P \equiv \sum_{k=1}^K \frac{1}{\sigma_k^2},$$

where  $\mathcal{N}(\mu, \sigma^2)$  is a normal distribution with mean  $\mu$  and variance  $\sigma^2$ ,  $Z_P$  is the precision of the models, and each individual model is assumed to possess a Gaussian form such as

$$\mathcal{M}_k \sim \mathcal{N}(f_k(x), \sigma_k^2(x)).$$

Here,  $f_k(x)$  is the mean of the model  $k$ , and  $\sigma_k^2(x)$  its variance, both at input parameter  $x$ .

In this method, the software receives the one-dimensional input space  $X$ , the mean of the  $k$  models at each point  $x \in X$  (hence it is a mean-based mixing procedure), and the variances of the models at each point  $x \in X$ . Each model is assumed to have been calibrated prior to being included in the mix. The ignorance of this mixing method with respect to how each model was generated allows for any model to be used, including Bayesian Machine Learning tools such as Gaussian Processes (Sempowski et al., 2022a) and Bayesian Neural Networks (Kronheim et al., 2022).

## Model mixing using Bayesian additive regression trees


A third BMM approach implemented in Taweret adopts a mean-mixing strategy which models the weight functions using Bayesian Additive Regression Trees (BART) conditional on the mean predictions from a set of  $K$  models (Yannotty et al., 2023). This approach enables the weight functions to be adaptively learned using tree bases and avoids the need for user-specified basis functions (such as a generalized linear model). Formally, the weight functions are defined by

$$w_k(x) = \sum_{j=1}^m g_k(x; T_j, M_j), \quad \text{for } k = 1, \dots, K$$

where  $g_k(x; T_j, M_j)$  defines the  $k^{\text{th}}$  output of the  $j^{\text{th}}$  tree,  $T_j$ , using the associated set of parameters,  $M_j$ . Each weight function is implicitly regularized via a prior to prefer the interval  $[0, 1]$ . Furthermore, the weight functions are not required to sum-to-one and can take values outside of the range of  $[0, 1]$ . This regularization approach is designed to maintain the flexibility of the model while also encouraging the weight functions to take values which preserve desired inferential properties.

This BART-based approach is implemented in C++ with the trees module in Taweret acting as a Python interface. The C++ back-end implements Bayesian tree models and originates from the *Open Bayesian Trees Project* (OpenBT) (Pratola et al., 2023). This module serves as an example for how existing code bases can be integrated into the Taweret framework.

## 105 Overview of package structure



**Figure 1:** Diagram depicting the base classes, their methods (functions) and their properties (data).

106 Taweret uses abstract base classes to ensure compatibility and uniformity of models and mixing  
 107 methods. The two base classes are **BaseModel** and **BaseMixer** located in the core folder (see  
 108 Fig. 1 for a schematic); any model mixing method developed with Taweret is required to  
 109 inherit from these. The former represents physics-based models that may include parameters  
 110 which need to be determined by Bayesian inference. The latter, **BaseMixer**, represents a mixing  
 111 method used to combine the predictions from the physics-based models using Bayesian model  
 112 mixing.

113 The design philosophy for Taweret is to make it easy to switch between mixing methods  
 114 without having to rewrite an analysis script. Thus, the base classes prescribe which functions  
 115 need to be present for interoperability between mixing methods, and in particular, the models  
 116 being called in the method. The functions required by **BaseModel** are

- 117 ■ **evaluate** - gives a point prediction for the model;
- 118 ■ **log\_likelihood\_elementwise** - calculates the log-likelihood, reducing along the last  
 119 axis of an array if the input array has multiple axes;
- 120 ■ **set\_prior** - sets priors for parameters in the model.

121 The functions required by **BaseMixer** are

- 122 ■ **evaluate** - gives point prediction for the mixed model given a set of parameters;
- 123 ■ **evaluate\_weights** - gives point prediction for the weights given a set of parameters;
- 124 ■ **map** - returns the maximum *a posteriori* estimate for the parameters of the mixed model  
 125 (which includes both the weights and any model parameters);

- 126     ▪ posterior - returns the chains of the sampled parameters from the mixed model;
- 127     ▪ predict - returns the posterior predictive distribution for the mixed model;
- 128     ▪ predict\_weights - returns the posterior predictive distribution for the model weights;
- 129     ▪ prior - returns the prior distributions (typically objects, not arrays);
- 130     ▪ prior\_predict - returns the prior predictive distribution for the mixed model;
- 131     ▪ set\_prior - sets the prior distributions for the mixing method;
- 132     ▪ train - executes the Bayesian model mixing step.

133     Following our design philosophy, the general workflow for an analysis using Taweret is described  
 134     in Fig. 2. From this, one can see three sources of information are generally required for an  
 135     analysis: a selected mixing method, a model set, and training data. Each of these sources are  
 136     connected through the training phase, which is where the mixing weights are learned. This  
 137     leads into the prediction phase, where final predictions are obtained for the overall system and  
 138     the weight functions. This process is summarized in the code snippet below. This workflow is  
 139     preserved across the various methods implemented in Taweret and is intended to be maintained  
 140     for future mixing methods included in this work.

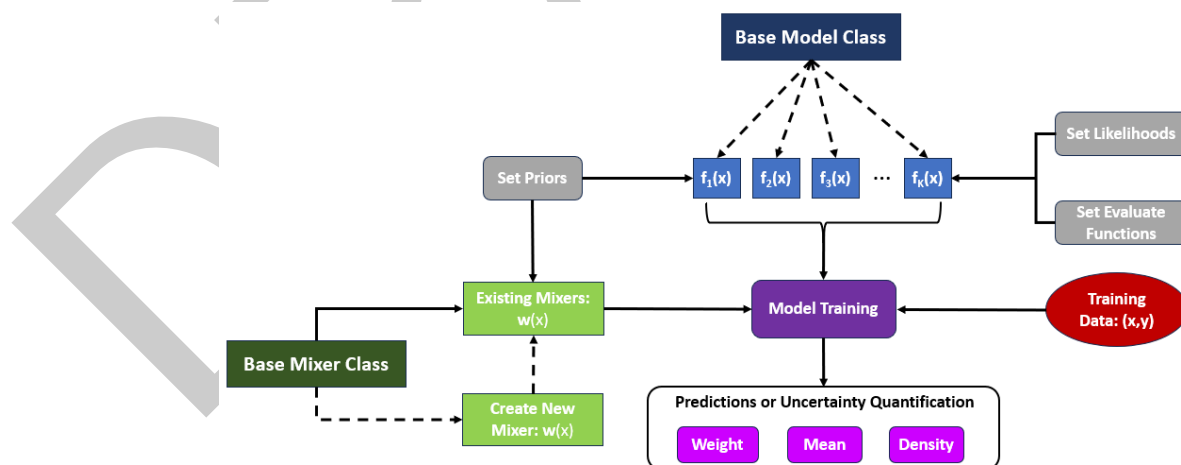
```

from mix.mix_method import MixMethod
from models.my_model import MyModel

mixer = MixMethod(models={'model_1': MyModel(...), ...})
mixer.set_prior(...)
mixer.train(...)
mixer.predict(...)
mixer.predict_weights(...)

```

141     Extending Taweret with a custom class or model simply requires that you inherit from the  
 142     base classes and implement the required functions.



**Figure 2:** The general workflow for an analysis using Taweret. (Blue) The user must define each of the  $K$  models as a class inherited from `BaseModel`. (Green) The user can select an existing mixing method from Taweret (solid) or contribute a new method (dashed). (Purple) The model is trained using a set of training data (red), the model set (blue), and the selected mixing method (green). Predictions and uncertainty quantification follow from the training process.

## Taweret moving forward

There are certainly many improvements that can be made to Taweret. An obvious one is a generalization of the bivariate linear mixing; this could be the mixing of an arbitrary number of models at the density level. Complementary to this density mixing method is a stochastic, mean-mixing method of arbitrary number  $K$  of models. An extension of the Multivariate Mixing method to multi-dimensional input and output spaces, correlated models, as well as calibration during mixing, is anticipated in future releases. Lastly, to facilitate the utilization of this growing framework, we hope to enable continuous integration routines for individuals contributing and create docker images that will run Taweret.

## Disclosure statement

The authors of this work are not aware of any conflicts of interest that would affect this research.

## Acknowledgments

We thank Daniel R. Phillips, Ulrich Heinz, Matt Pratola, Kyle Godbey, Stefan Wild, Sunil Jaiswal, and all other BAND members for crucial feedback and discussion during the development stage of this package. This work is supported by the CSSI program Award OAC-2004601 (DL, ACS, JCY). ACS also acknowledges support from the Department of Energy (contract no. DE-FG02-93ER40756).

## References

- Beyer, K., Buskirk, L., Chan, M. Y.-H., Chang, T. H., DeBoer, R. J., Furnstahl, R. J., Giuliani, P., Godbey, K., Ingles, K., Liyanage, D., Nunes, F. M., Odell, D., Phillips, D. R., Plumlee, M., Pratola, M. T., Sempowski, A. C., Sürer, Ö., Wild, S. M., & Yannotty, J. C. (2023). *BANDFramework: An open-source framework for Bayesian analysis of nuclear dynamics* (Version 0.3.0). <https://github.com/bandframework/bandframework>
- FitzGerald, T. H. B., Dolan, R. J., & Friston, K. J. (2014). Model averaging, optimal inference, and habit formation. *Frontiers in Human Neuroscience*, 8:457. <https://doi.org/10.3389/fnhum.2014.00457>
- Fragoso, T. M., Bertoli, W., & Louzada, F. (2018). Bayesian model averaging: A systematic review and conceptual classification. *International Statistical Review*, 86(1), 1–28. <https://doi.org/https://doi.org/10.1111/insr.12243>
- Kronheim, B., Kuchera, M., & Prosper, H. (2022). TensorBNN: Bayesian inference for neural networks using TensorFlow. *Computer Physics Communications*, 270, 108168. <https://doi.org/10.1016/j.cpc.2021.108168>
- Liyanage, D. (2023). *Multifaceted study of ultrarelativistic heavy ion collisions* [PhD thesis, Ohio State University; The Ohio State University]. [http://rave.ohiolink.edu/etdc/view?acc\\_num=osu1689508609203123](http://rave.ohiolink.edu/etdc/view?acc_num=osu1689508609203123)
- Phillips, D. R., Furnstahl, R. J., Heinz, U., Maiti, T., Nazarewicz, W., Nunes, F. M., Plumlee, M., Pratola, M. T., Pratt, S., Viens, F. G., & Wild, S. M. (2021). Get on the BAND Wagon: A Bayesian Framework for Quantifying Model Uncertainties in Nuclear Dynamics. *Journal of Physics G*, 48(7), 072001. <https://doi.org/10.1088/1361-6471/abf1df>
- Pratola, M. T., McCulloch, R. E., Chipman, H. A., & Horiguchi, A. (2023). *Open Bayesian Trees Project*. <https://bitbucket.org/mpratola/openbt/wiki/Home>

- 185 Raftery, A., Hoeting, J., Volinsky, C., Painter, I., & Yeung, K. Y. (2022). *BMA: Bayesian*  
186 *model averaging*. <https://CRAN.R-project.org/package=BMA>
- 187 Semposki, A. C., Furnstahl, R. J., & Phillips, D. R. (2022a). Interpolating between small-  
188 and large-g expansions using Bayesian model mixing. *Physical Review C*, 106(4), 044002.  
189 <https://doi.org/10.1103/PhysRevC.106.044002>
- 190 Semposki, A. C., Furnstahl, R. J., & Phillips, D. R. (2022b). *SAMBA: SAndbox for Mixing*  
191 *using Bayesian Analysis*. <https://github.com/aseposki/SAMBA>
- 192 Sloughter, J. M., Raftery, A. E., Gneiting, T., & Fraley, C. (2007). Probabilistic quantitative  
193 precipitation forecasting using Bayesian model averaging. *Monthly Weather Review*, 135,  
194 3209–3220. <https://doi.org/10.1175/MWR3441.1>
- 195 Vehtari, A., Gelman, A., & Gabry, J. (2017). *loo: Efficient leave-one-out cross-validation and*  
196 *WAIC for Bayesian models*. <https://mc-stan.org/loo/>
- 197 Yannoty, J. C., Santner, T. J., Furnstahl, R. J., & Pratola, M. T. (2023). Model mixing  
198 using Bayesian additive regression trees. In *Technometrics* (pp. 1–12). Taylor & Francis.  
199 <https://doi.org/10.1080/00401706.2023.2257765>
- 200 Yao, Y., Pirš, G., Vehtari, A., & Gelman, A. (2022). Bayesian Hierarchical Stacking: Some  
201 Models Are (Somewhere) Useful. *Bayesian Analysis*, 17(4), 1043–1071. <https://doi.org/10.1214/21-BA1287>  
202
- 203 Yao, Y., Vehtari, A., Simpson, D., & Gelman, A. (2018). Using stacking to average bayesian  
204 predictive distributions (with discussion). *Bayesian Analysis*, 13 (3) 917 - 1007. <https://doi.org/10.1214/17-BA1091>  
205