

Surjectors: surjection layers for density estimation with normalizing flows

Simon Dirmeier^{1,2}

¹ Swiss Data Science Center, Zurich, Switzerland ² ETH Zurich, Zurich, Switzerland

DOI: [10.xxxxxx/draft](https://doi.org/10.xxxxxx/draft)

Software

- [Review](#) ↗
- [Repository](#) ↗
- [Archive](#) ↗

Editor: [Arfon Smith](#) ↗

Reviewers:

- [@sandeshkatakam](#)
- [@animikhaich](#)

Submitted: 14 October 2023

Published: unpublished

License

Authors of papers retain copyright and release the work under a Creative Commons Attribution 4.0 International License ([CC BY 4.0](#)).

Summary

Normalizing flows (NFs, [Papamakarios et al., 2021](#)) are tractable neural density estimators which have in the recent past been applied successfully for, e.g., generative modelling [Ping et al. \(2020\)](#), Bayesian inference [Hoffman et al. \(2019\)](#) or simulation-based inference [Dirmeier et al. \(2023\)](#). Surjectors is a Python library in particular for *surjective*, i.e., dimensionality-reducing normalizing flows (SNFs, [Klein et al. \(2021\)](#)). Surjectors is based on the libraries JAX, Haiku and Distrax [Babuschkin et al. \(2020\)](#) and is fully compatible with them. By virtue of being entirely written in JAX ([Bradbury et al., 2018](#)), Surjectors naturally supports usage on either CPU, GPU or TPU.

Statement of Need

Real-world data are often lying in a high-dimensional ambient space embedded in a lower-dimensional manifold ([Fefferman et al., 2016](#)) which can complicate estimation of probability densities ([Dai & Seljak \(2021\)](#), [Klein et al. \(2021\)](#), [Nalisnick et al. \(2019\)](#)). As a remedy, recently neural density estimators using surjective normalizing flows (SNFs) have been proposed which reduce the dimensionality of the data while still allowing for exact computation of data likelihoods ([Klein et al., 2021](#)). While several computational libraries exist that implement *bijective* normalizing flows, i.e., flows that are dimensionality-preserving, currently none exist that efficiently implement dimensionality-reducing flows.

Surjectors is a normalizing flow library that implements both bijective and surjective normalizing flows. Surjectors is light-weight, conceptually simple to understand if familiar with the JAX ecosystem, and computationally efficient due to leveraging the XLA compilation and vectorization from JAX. We additionally make use of several well-established packages within the JAX ecosystem ([Bradbury et al., 2018](#)) and probabilistic deep learning community. For composing the conditioning networks that NFs facilitate, Surjectors uses the deep learning library Haiku ([Hennigan et al., 2020](#)). For training and optimisation, we utilize the gradient transformation library Optax ([Babuschkin et al., 2020](#)). Surjectors leverages Distrax ([Babuschkin et al., 2020](#)) and TensorFlow probability ([Dillon et al., 2017](#)) for probability distributions and several base bijector implementations.

Adoption

[Dirmeier et al. \(2023\)](#) have proposed a novel method for simulation-based inference where they make use autoregressive inference surjections for density estimation and where they are using Surjectors for their implementations.

References

- 37
- 38 Babuschkin, I., Baumli, K., Bell, A., Bhupatiraju, S., Bruce, J., Buchlovsky, P., Budden,
39 D., Cai, T., Clark, A., Danihelka, I., Dedieu, A., Fantacci, C., & others. (2020). *The*
40 *DeepMind JAX Ecosystem*. <http://github.com/deepmind>
- 41 Bradbury, J., Frostig, R., Hawkins, P., Johnson, M. J., Leary, C., Maclaurin, D., Necula, G.,
42 Paszke, A., VanderPlas, J., Wanderman-Milne, S., & Zhang, Q. (2018). *JAX: Composable*
43 *transformations of Python+NumPy programs*. <http://github.com/google/jax>
- 44 Dai, B., & Seljak, U. (2021). Sliced iterative normalizing flows. *ICML Workshop on Invertible*
45 *Neural Networks, Normalizing Flows, and Explicit Likelihood Models*.
- 46 Dillon, J. V., Langmore, I., Tran, D., Brevdo, E., Vasudevan, S., Moore, D., Patton, B.,
47 Alemi, A., Hoffman, M., & Saurous, R. A. (2017). Tensorflow distributions. *arXiv Preprint*
48 *arXiv:1711.10604*.
- 49 Dirmeier, S., Albert, C., & Perez-Cruz, F. (2023). Simulation-based inference using surjective
50 sequential neural likelihood estimation. *arXiv Preprint arXiv:2308.01054*.
- 51 Fefferman, C., Mitter, S., & Narayanan, H. (2016). Testing the manifold hypothesis. *Journal*
52 *of the American Mathematical Society*, 29(4), 983–1049.
- 53 Hennigan, T., Cai, T., Norman, T., Martens, L., & Babuschkin, I. (2020). *Haiku: Sonnet for*
54 *JAX*. <http://github.com/deepmind/dm-haiku>
- 55 Hoffman, M., Sountsov, P., Dillon, J. V., Langmore, I., Tran, D., & Vasudevan, S. (2019).
56 Neutra-lizing bad geometry in hamiltonian monte carlo using neural transport. *arXiv*
57 *Preprint arXiv:1903.03704*.
- 58 Kingma, D. P., & Dhariwal, P. (2018). Glow: Generative flow with invertible 1x1 convolutions.
59 *Advances in Neural Information Processing Systems*.
- 60 Klein, S., Raine, J. A., Pina-Otey, S., Voloshynovskiy, S., & Golling, T. (2021). Funnels: Exact
61 maximum likelihood with dimensionality reduction. *Workshop on Bayesian Deep Learning,*
62 *Advances in Neural Information Processing Systems*.
- 63 Nalisnick, E., Matsukawa, A., Teh, Y. W., Gorur, D., & Lakshminarayanan, B. (2019). Do
64 deep generative models know what they don't know? *International Conference on Learning*
65 *Representations*.
- 66 Papamakarios, G., Nalisnick, E., Rezende, D. J., Mohamed, S., & Lakshminarayanan, B.
67 (2021). Normalizing flows for probabilistic modeling and inference. *The Journal of Machine*
68 *Learning Research*, 22(1), 2617–2680.
- 69 Papamakarios, G., Sterratt, D., & Murray, I. (2019). Sequential neural likelihood: Fast
70 likelihood-free inference with autoregressive flows. *Proceedings of the 22nd International*
71 *Conference on Artificial Intelligence and Statistics*.
- 72 Ping, W., Peng, K., Zhao, K., & Song, Z. (2020). WaveFlow: A compact flow-based model
73 for raw audio. *Proceedings of the 37th International Conference on Machine Learning*.
- 74 Rezende, D., & Mohamed, S. (2015). Variational inference with normalizing flows. *Proceedings*
75 *of the 32nd International Conference on Machine Learning*.